| Wednesday 04/12 | |
|---|---|
| **Main Room, 16h30** | |
| Archaeoscape: Bringing Aerial Laser Scanning Archaeology to the Deep Learning Era | Yohann Perron |
| Towards training digitally-tied analog blocks via hybrid gradient computation | Maxence Ernoult |
| WFCRL: A Multi-Agent Reinforcement Learning Benchmark for Wind Farm Control | Claire Bizon-Monroc |
| Binding in hippocampal-entorhinal circuits enables compositionality in cognitive maps | Sonia Mazelet |
| A generalized neural tangent kernel for surrogate gradient learning | Luke Eilers |
| Dimension-free deterministic equivalents for random feature regression | Leonardo Defilippis |
| Barely Random Algorithms for Metrical Task Systems | Romain Cosson |
| Statistical and Geometrical properties of Kernel Kullback-Leibler divergence | Clémentine Chazal |
| Topological Generalization Bounds for Discrete-Time Stochastic Optimization Algorithms | Benjamin Dupuis |
| Optimal Classification under Performative Distribution Shift | Edwige Cyffers, Olivier Cappé, JamalAtif |
| Nonconvex Federated Learning on Compact Smooth Submanifolds With Heterogeneous | Jiaojiao Zhang |
| Non-asymptotic Analysis of Biased Adaptive Stochastic Approximation | Adeline Fermanian, Sohiban Surendran |
| A Novel Approach to Loss Landscape Characterization without Over-Parametrization | Rustem Islamov |
| Variational Graph Contrastive Learning | shifeng xie |
| In-context Quantile Regression for Multi-product Inventory Management using Time-seri | Sohom Mukherjee |
| Bandits with Abstention under Expert Advice | Maximilian Thiessen |
| An Analysis of Elo Rating Systems via Markov Chains | Luca Zanetti |
| SCAFFLSA: Taming Heterogeneity in Federated Linear Stochastic Approximation and TI | Paul Mangold |
| DEFT: Efficient Finetuning of Conditional Diffusion Models by Learning the Generalised $h$-transform | Alexander Denker |
| A Unifying Post-Processing Framework for Multi-Objective Learn-to-Defer Problems | Amin Charusaie |
| Semi-Discrete Optimal Transport: Nearly Minimax Estimation With Stochastic Gradient D | Ferdinand Genans, Antoine Godichon-Baggioni |
| An eye for an ear: zero-shot audio description leveraging an image captioner with audio- | Hugo Malard |
| When is an Embedding Model More Promising than Another | Maxime Darrin |
| Metacognitive Capabilities of LLMs: An Exploration in Mathematical Problem Solving | Michal Valko |
| Boosting Generalization in Parametric PDE Neural Solvers through Adaptive Conditionin | ArmandKassai |
| Diffeomorphic interpolation for efficient persistence-based topological optimization | Théo Lacombe |
| 6485, 15804, 17579 | Marco Cuturi |
| | |
| **Secondary Room, 16h30** | |
| Learning to Mitigate Externalities: the Coase Theorem with Hindsight Rationality | Antoine Scheid |
| Logarithmic Smoothing for Pessimistic Off-Policy Evaluation, Selection and Learning | Imad Aouali, Otmane Sakhi |
| Extensive-Form Game Solving via Blackwell Approachability on Treeplexes.<br><br>Fast Last-Iterate Convergence of Learning in Games Requires Forgetful Algorithms. | Julien Grand-Clément |
| The Value of Reward Lookahead in Reinforcement Learning<br>Reinforcement Learning with Lookahead Information<br>Improved Algorithms for Contextual Dynamic Pricing | Nadav Merlis |
| MetaCURL: Non-stationary Concave Utility Reinforcement Learning | Bianca Marin Moreno |
| A Concept-Based Explainability Framework for Large Multimodal Models | Pegah KHAYATAN, Jayneel Parekh |
| General Detection-based Text Line Recognition | Syrine Kalleli, Raphael Baena |
| Almost Free: Self-concordance in Natural Exponential Families and an Application to Ba | Flore Sentenac |
| Bridging semantics and pragmatics in information-theoretic emergent communication | Eleonora Gualdoni |
| DEFT: Efficient Finetuning of Conditional Diffusion Models by Learning the Generalised -transform<br>Improving Linear System Solvers for Hyperparameter Optimisation in Iterative Gaussian | ShreyasPadhy |
| Only Strict Saddles in the Energy Landscape of Predictive Coding Networks? | El MehdiAchour |
| Combining Statistical Depth and Fermat Distance for Uncertainty Quantification | Hai-Vy NGUYEN, RedaCHHAIBI |
| Causal Contrastive Learning for Counterfactual Regression Over Time | MouadEL Bouchattaoui |
| MaNo: Exploiting Matrix Norm for Unsupervised Accuracy Estimation under Distribution | Ambroise Odonnat |
| The Well: a Large-Scale Collection of Diverse Physics Simulations for Machine Learning | Lucas Meyer |
| Shape analysis for time series | Samuel Gruffaz |
| Confidence Calibration of Classifiers with Many Classes | Adrien Le Coz |
| Iteration heads | Vivien Cabannes |
| I Don't Know: Explicit Modeling of Uncertainty with an [IDK] Token | Konstantin Dobler |
| MicroAdam: Accurate Adaptive Optimization with Low Space Overhead and Provable Co | ThomasROBERT |
| Model-free Low-Rank Reinforcement Learning via Leveraged Entry-wise Matrix Estimati | Alexandre Proutiere |
| Unravelling in Collaborative Learning | Aymeric Capitaine |
| DU-Shapley: A Shapley Value Proxy for Efficient Dataset Valuation | Maxime Vono |
| DiffCut: Catalyzing Zero-Shot Semantic Segmentation with Diffusion Features and Recu | PaulCouairon |

| | |
|---|---|
| Near-Optimal Distributionally Robust RL with General $L_p$ Norms<br><br>Time-Constrained Robust MDPs | <br><br>Pierre Clavier |
| | |
| | |
| **Thursday 05/12** | |
| **Main Room, 12h30** | |
| Watermarking Makes Language Models Radioactive | Pierre Fernandez, Tom Sander |
| Benchmarking Uncertainty Disentanglement: Specialized Uncertainties for Specialized T | Michael Kirchhof |
| FairJob: A Real-World Dataset for Fairness in Online Systems | Mariia Vladimirova |
| Consent in Crisis: The Rapid Decline of the AI Data Commons | Christopher Klamm |
| Functional Bilevel Optimization for Machine Learning | Ieva Petrulionyte, Julien Mairal |
| Mirror and Preconditioned Gradient Descent in Wasserstein Space | Clément Bonet |
| The Road Less Scheduled | Konstantin Mishchenko |
| What makes unlearning hard and what to do about it | Kairan Zhao |
| Learning with Fitzpatrick Losses | Seta Rakotomandimby, Michel De Lara, Mathieu Blondel |
| Learning to Embed Distributions via Maximum Kernel Entropy | Oleksii Kachaiev |
| Piecewise deterministic generative models | Dario Shariatian |
| Annealed Multiple Choice Learning: Overcoming limitations of Winner-takes-all with ann | David Perera |
| Implicit Bias of Mirror Flow on Separable Data | Scott Pesme, Radu Dragomir |
| Learning the Infinitesimal Generator of Stochastic Diffusion Processes<br><br>From Biased to Unbiased Dynamics: An Infinitesimal Generator Approach | <br><br>VladimirKostic |
| Neural Conditional Probability for Inference | Vladimir Kostic, Karim Lounici |
| Expected Probabilistic Hierarchies<br>Shaving Weights with Occam's Razor: Bayesian Sparsification for Neural Networks usin | <br>BertrandCharpentier |
| Theoretical guarantees in KL for Diffusion Flow Matching | Alain Oliviero-Durmus |
| Near-Optimality of Contrastive Divergence Algorithms | Pierre Glaser |
| | |
| **Main Room, 16h30** | |
| Any2Graph: Deep End-To-End Supervised Graph Prediction With An Optimal Transport | Paul Krzakala, Rémi Flamary, Florence d'Alché-Buc |
| Analysing Multi-Task Regression via Random Matrix Theory with Application to Time Ser | Vasilii Feofanov |
| ANAH-v2: Scaling Analytical Hallucination Annotation of Large Language Models | Ziwei Ji |
| Understanding Visual Feature Reliance through the Lens of Complexity | Louis Bethune |
| Towards Efficient and Optimal Covariance-Adaptive Algorithms for Combinatorial Semi-E | Julien Zhou, Thibaud Rahier |
| Supra-Laplacian Encoding for Transformer on Dynamic Graphs | Yannis Karmim |
| ManiPose: Manifold-Constrained Multi-Hypothesis 3D Human Pose Estimation | Victor Letzelter |
| Continuous Product Graph Neural Networks | Aref EINIZADE, Jhony H. Giraldo |
| Wormhole loss for partial shape matching | Thomas Dagès |
| Improved learning rates in multi-unit uniform price auctions | Hugo Richard, MariusPotfer |
| Optimizing the coalition gain in Online Auctions with Greedy Structured Bandits | Hugo Richard, Dorian Baudry |
| Deep linear networks for regression are implicitly regularized towards flat minima | Pierre Marion |
| BOLD: Boolean Logic Deep Learning | Van Minh Nguyen, Ba-HienTran |
| AROMA: Preserving Spatial Structure for Latent PDE Modeling with Local Neural Field | Louis Serrano, Jean-Noel Vittaut |
| Implicit Multimodal Alignment: On the Generalization of Frozen LLMs to Multimodal Inpu | MustafaShukor |
| You Don't Need Data-Augmentations in Self-Supervised Learning | Théo Moutakanni |
| Aligning Embeddings and Geometric Random Graphs: Informational Results and Compu | Mathieu Even, Luca Ganassali, Jakob Maier |
| Overcoming Brittleness in Pareto-Optimal Learning Augmented Algorithms | Christoph Dürr |
| Computing the Bias of Constant-step Stochastic Approximation with Markovian Noise | Nicolas Gast |
| Activation Map Compression through Tensor Decomposition for Deep Learning | Enzo Tartaglione, AëlQuélennec |
| Improving Neural Network Surface Processing with Principal Curvatures | JosquinHarrison |
| Divide-and-Conquer Posterior Sampling for Denoising Diffusion priors | Badr Moufad |
| Credal Learning Theory | Michele Caprio |