Wednesday 04/12	
Caves d'Esclangon (floor -1), 16h30	
Learning to Mitigate Externalities: the Coase Theorem with Hindsight Rationality	Antoine Scheid
Logarithmic Smoothing for Pessimistic Off-Policy Evaluation, Selection and Learning	Imad Aouali, Otmane Sakhi
Extensive-Form Game Solving via Blackwell Approachability on Treeplexes.	Julien Grand-Clément
Fast Last-Iterate Convergence of Learning in Games Requires Forgetful Algorithms.	
The Value of Reward Lookahead in Reinforcement Learning	Nadav Merlis
Reinforcement Learning with Lookahead Information	
Improved Algorithms for Contextual Dynamic Pricing	
MetaCURL: Non-stationary Concave Utility Reinforcement Learning	Bianca Marin Moreno
A Concept-Based Explainability Framework for Large Multimodal Models	Pegah Khayatan, Jayneel Parekh
Almost Free: Self-concordance in Natural Exponential Families and an Application to Bandits	Flore Sentenac
DEFT: Efficient Finetuning of Conditional Diffusion Models by Learning the Generalised	Shreyas Padhy, Alexander Denker
Causal Contrastive Learning for Counterfactual Regression Over Time	Mouad El Bouchattaoui
Shape analysis for time series	Samuel Gruffaz
Confidence Calibration of Classifiers with Many Classes	Adrien Le Coz
Model-free Low-Rank Reinforcement Learning via Leveraged Entry-wise Matrix Estimation	Alexandre Proutiere
Unravelling in Collaborative Learning	Aymeric Capitaine
DU-Shapley: A Shapley Value Proxy for Efficient Dataset Valuation	Maxime Vono
Near-Optimal Distributionally Robust RL with General Lp Norms	Pierre Clavier
Time-Constrained Robust MDPs	
SCAI (floor 1), 16h30	
A generalized neural tangent kernel for surrogate gradient learning	Luke Eilers
Dimension-free deterministic equivalents for random feature regression	Leonardo Defilippis
Barely Random Algorithms for Metrical Task Systems	Romain Cosson
Statistical and Geometrical properties of Kernel Kullback-Leibler divergence	Clémentine Chazal
Topological Generalization Bounds for Discrete-Time Stochastic Optimization Algorithms	Benjamin Dupuis
Optimal Classification under Performative Distribution Shift	Edwige Cyffers, Olivier Cappé, Jamal Atif
Nonconvex Federated Learning on Compact Smooth Submanifolds With Heterogeneous Data	Jiaojiao Zhang
Non-asymptotic Analysis of Biased Adaptive Stochastic Approximation	Adeline Fermanian, Sohiban Surendran, Antoine Godichon-Baggioni
A Novel Approach to Loss Landscape Characterization without Over-Parametrization	Rustem Islamov
Variational Graph Contrastive Learning	Shifeng Xie
In-context Quantile Regression for Multi-product Inventory Management using Time-series Transformers	Sohom Mukherjee
Bandits with Abstention under Expert Advice	Maximilian Thiessen
An Analysis of Elo Rating Systems via Markov Chains	Luca Zanetti
SCAFFLSA: Taming Heterogeneity in Federated Linear Stochastic Approximation and TD Learning	Paul Mangold
A Unifying Post-Processing Framework for Multi-Objective Learn-to-Defer Problems	Amin Charusaie
Semi-Discrete Optimal Transport: Nearly Minimax Estimation With Stochastic Gradient Descent and Adaptive Entropic Regularization	Ferdinand Genans, Antoine Godichon-Baggioni
Metacognitive Capabilities of LLMs: An Exploration in Mathematical Problem Solving	Michal Valko
Diffeomorphic interpolation for efficient persistence-based topological optimization	Théo Lacombe
Progressive Entropic Optimal Transport Solvers Parnian Kassraie	Marco Cuturi
Learning Elastic Costs to Shape Monge Displacements	
GENOT: A Neural Optimal Transport Framework for Single Cell Genomics	
Hall d'Esclangon (floor 0), 16h30	
Archaeoscape: Bringing Aerial Laser Scanning Archaeology to the Deep Learning Era	Yohann Perron
Towards training digitally-tied analog blocks via hybrid gradient computation	Maxence Ernoult
WFCRL: A Multi-Agent Reinforcement Learning Benchmark for Wind Farm Control	Claire Bizon-Monroc
Binding in hippocampal-entorhinal circuits enables compositionality in cognitive maps	Sonia Mazelet
An eye for an ear: zero-shot audio description leveraging an image captioner with audio-visual token distribution matching	Hugo Malard
When is an Embedding Model More Promising than Another	Maxime Darrin
Boosting Generalization in Parametric PDE Neural Solvers through Adaptive Conditioning	Armand Kassai, Jean-Noel Vittaut
General Detection-based Text Line Recognition	Syrine Kalleli, Raphael Baena
Bridging semantics and pragmatics in information-theoretic emergent communication	Eleonora Gualdoni
Improving Linear System Solvers for Hyperparameter Optimisation in Iterative Gaussian Processes	Shreyas Padhy
Only Strict Saddles in the Energy Landscape of Predictive Coding Networks?	El Mehdi Achour

Combining Statistical Depth and Fermat Distance for Uncertainty Quantification	Hai-Vy Nguyen, Reda Chhaibi
The Well: a Large-Scale Collection of Diverse Physics Simulations for Machine Learning	Lucas Meyer
Iteration heads: A Mechanistic Study of Chain-of-Thought	Vivien Cabannes
MicroAdam: Accurate Adaptive Optimization with Low Space Overhead and Provable Convergence	Thomas Robert
DiffCut: Catalyzing Zero-Shot Semantic Segmentation with Diffusion Features and Recursive Normalized Cut	Paul Couairon
MaNo: Exploiting Matrix Norm for Unsupervised Accuracy Estimation under Distribution Shifts	Ambroise Odonnat, Vasilii Feofanov
I Don't Know: Explicit Modeling of Uncertainty with an [IDK] Token	Konstantin Dobler

Thursday 05/12

Hall d'Esclangon (floor 0), 12h30

Watermarking Makes Language Models Radioactive	Pierre Fernandez, Tom Sander
Benchmarking Uncertainty Disentanglement: Specialized Uncertainties for Specialized Tasks	Michael Kirchhof
FairJob: A Real-World Dataset for Fairness in Online Systems	Mariia Vladimirova
Consent in Crisis: The Rapid Decline of the AI Data Commons	Christopher Klamm
Functional Bilevel Optimization for Machine Learning	Ieva Petrulionyte, Julien Mairal
Mirror and Preconditioned Gradient Descent in Wasserstein Space	Clément Bonet
The Road Less Scheduled	Konstantin Mishchenko
What makes unlearning hard and what to do about it	Kairan Zhao
Learning with Fitzpatrick Losses	Seta Rakotomandimby, Michel De Lara, Mathieu Blondel
Learning to Embed Distributions via Maximum Kernel Entropy	Oleksii Kachaiev
Piecewise deterministic generative models	Dario Shariatian
Annealed Multiple Choice Learning: Overcoming limitations of Winner-takes-all with annealing,	David Perera
Implicit Bias of Mirror Flow on Separable Data	Scott Pesme, Radu Dragomir
Learning the Infinitesimal Generator of Stochastic Diffusion Processes	Vladimir Kostic
From Biased to Unbiased Dynamics: An Infinitesimal Generator Approach	
Neural Conditional Probability for Inference	Vladimir Kostic, Karim Lounici
Expected Probabilistic Hierarchies	Bertrand Charpentier
Shaving Weights with Occam's Razor: Bayesian Sparsification for Neural Networks using the Marginal Likelihood	
Theoretical guarantees in KL for Diffusion Flow Matching	Alain Oliviero-Durmus, Marta Gentiloni Silveri
Near-Optimality of Contrastive Divergence Algorithms	Pierre Glaser

SCAI (floor 1), 12h30

Any2Graph: Deep End-To-End Supervised Graph Prediction With An Optimal Transport Loss	Paul Krzakala, Rémi Flamary, Florence d'Alché-Buc
Analysing Multi-Task Regression via Random Matrix Theory with Application to Time Series Forecasting	Vasilii Feofanov
ANAH-v2: Scaling Analytical Hallucination Annotation of Large Language Models	Ziwei Ji
Understanding Visual Feature Reliance through the Lens of Complexity	Louis Bethune
Towards Efficient and Optimal Covariance-Adaptive Algorithms for Combinatorial Semi-Bandits	Julien Zhou, Thibaud Rahier
Supra-Laplacian Encoding for Transformer on Dynamic Graphs	Yannis Karmim
ManiPose: Manifold-Constrained Multi-Hypothesis 3D Human Pose Estimation	Victor Letzelter
Continuous Product Graph Neural Networks	Aref Einizade, Jhony H. Giraldo
Wormhole loss for partial shape matching	Thomas Dagès
Improved learning rates in multi-unit uniform price auctions	Hugo Richard, Marius Potfer
Optimizing the coalition gain in Online Auctions with Greedy Structured Bandits	Hugo Richard, Dorian Baudry
Deep linear networks for regression are implicitly regularized towards flat minima	Pierre Marion
BOLD: Boolean Logic Deep Learning	Van Minh Nguyen, Ba-Hien Tran
AROMA: Preserving Spatial Structure for Latent PDE Modeling with Local Neural Field	Louis Serrano, Jean-Noel Vittaut
Implicit Multimodal Alignment: On the Generalization of Frozen LLMs to Multimodal Inputs	Mustafa Shukor
You Don't Need Data-Augmentations in Self-Supervised Learning	Théo Moutakanni
Aligning Embeddings and Geometric Random Graphs: Informational Results and Computational Approaches for the Procrustes-Wasserstein Problem	Mathieu Even, Luca Ganassali, Jakob Maier
Overcoming Brittleness in Pareto-Optimal Learning Augmented Algorithms	Christoph Dürr
Computing the Bias of Constant-step Stochastic Approximation with Markovian Noise	Nicolas Gast
Activation Map Compression through Tensor Decomposition for Deep Learning	Enzo Tartaglione, Aël Quélennec
Improving Neural Network Surface Processing with Principal Curvatures	Josquin Harrison
Divide-and-Conquer Posterior Sampling for Denoising Diffusion priors	Badr Moufad
Credal Learning Theory	Michele Caprio